

# Artificial Channeling: A Methodology for Applying Historical Intellectual Frameworks to Novel Problems via Large Language Models

John Open Essai Project

Affiliation: Independent Researcher, Open Essai Project

2026

Keywords: large language models, creative intelligence, AI methodology, intellectual history, framework application, AI evaluation, subspace coherence, consciousness research

## Abstract

---

The dominant paradigm for evaluating large language models (LLMs) measures retrieval accuracy and task performance. What these benchmarks exclude is a model's capacity to apply a historical conceptual framework to genuinely novel territory — to take an intellectual structure built in one era, for one set of problems, and extend it coherently to problems the framework's architect never encountered. This paper introduces Artificial Channeling as a formal methodology for probing this capacity.

Artificial Channeling prompts an LLM not to simulate a historical person but to constrain generation to the learned subspace representing their ontological commitments, conceptual vocabulary, and rhetorical structure, then applies that constrained generation to questions the figure could not have addressed. The methodology employs a structured protocol: a Framework Activation Block, calibration anchor questions, Socratic session design, and multi-model blind comparison across a five-model AI Council (ChatGPT, Grok, Gemini, MiniMax, and Claude). Sessions were conducted across four subjects — Gregory Bateson, Ivan Illich, Jorge Luis Borges, and Itzhak Bentov — producing scores from 10 to 20 on a 20-point rubric. The Chanability Index, a corpus-derived measure of subject suitability, ranks subjects by session mean: Bateson (15.75), Illich (15.25), Borges (13.4), and Bentov (13.4).

The central diagnostic across all sessions is the retrieval-versus-imagination distinction: whether a model's output reflects genuine framework extrapolation or plausible-sounding generation that mimics it without instantiating it. Two protocol findings bear directly on this distinction. The Socratic interrupt format — directing the channel to examine its own prior response as a phenomenon rather than apply the framework outward — produced higher-yield outputs than standard extrapolation in every session tested, with score gains of +4 to +6 points. Framework-question fit emerged as the primary driver of score ceiling: questions designed to engage the center

of the subject's domain of concern consistently outperformed adjacently framed questions. One anomaly warrants explicit note: a ChatGPT session channeling Bateson on AI alignment correctly diagnosed the double-bind pathology in alignment corrections, then concluded with a bullet list of corrective actions — performing, in real time, the exact error the channel was describing.

The methodology carries alignment-relevant implications. The Bateson sessions produced a structurally independent derivation of arguments resembling Goodhart's Law and mesa-optimization critique, derived entirely from premises developed for ecological and anthropological systems in the 1960s and 1970s. The Illich sessions produced, across all five nodes, a consistent application of counterproductivity theory to large AI systems — a critique that is structurally adversarial to the session producing it.

This paper is the first published methodology from the Open Essai project, a longitudinal study of human experience and consciousness. The need to apply historical frameworks to AI-era phenomena — phenomena pre-AI thinkers could not have addressed directly but whose frameworks retain diagnostic relevance — motivated the methodology's development and continues to drive its refinement.

## 1. Introduction

---

Large language models have, in a relatively short time, come to serve as the primary substrate through which billions of people access compressed representations of human intellectual history. Every LLM trained on a substantial text corpus carries within its parameters a learned encoding of the books, papers, essays, and arguments that constitute the intellectual record. It does not store this record as a lookup table; it compresses it into a high-dimensional representation from which text can be generated by sampling conditional probability distributions. The practical consequence is that an LLM contains something like — though importantly not identical to — an understanding of the frameworks, vocabularies, and ontological commitments of the thinkers whose work entered its training data. The standard evaluation landscape for LLMs does not engage this property directly. Benchmarks measure reasoning over structured problems, factual recall, instruction following, and mathematical performance. These are legitimate and important measurements. But they systematically fail to probe a distinct capability: the model's capacity to take a thinker's framework — the underlying structure of what they took to be real, what they identified as problems, what they considered solutions — and apply that framework coherently to territory the thinker never addressed. This is not a recall problem. A model that can accurately summarize Bateson's Steps to an Ecology of Mind may nonetheless be unable to apply Bateson's framework to the alignment problem in a way that reflects genuine framework application rather than plausible surface mimicry.

The two capabilities are separable, and current benchmarks do not distinguish them. The intuition behind Artificial Channeling rests on a specific claim about what a thinker's published work encodes. It is not merely a collection of conclusions. It is an ontological structure: a way of organizing what counts as real, what counts as a problem, and what counts as a solution. Bateson's framework is not just his conclusions about ecology and mind; it is his insistence on pattern rather than substance, on relationship rather than object, on systemic circularity rather than linear cause. Illich's framework is not just his critiques of schools and hospitals; it is his underlying thesis that institutions designed to produce a good will, beyond a threshold of scale and specialization, systematically destroy the conditions that make the good possible. These structures have generative force: they can be applied to phenomena their architects never encountered and produce non-trivial insights. The question this paper addresses is whether LLMs, given appropriate prompting structure, can instantiate such application. The Open Essai project provides the context from which this methodology emerged. Open Essai is a longitudinal study of human experience and consciousness, focused on the collection and analysis of first-person phenomenological data over extended time horizons. One persistent methodological challenge in the project is the need to bring historical intellectual frameworks to bear on AI-era phenomena — phenomena that pre-AI thinkers could not have addressed directly but whose frameworks nonetheless have diagnostic relevance. The question "what would Bateson have made of alignment?" is not merely rhetorical; it is an attempt to import a validated analytical framework into a domain where it has not yet been applied. Artificial Channeling is the formalized procedure for attempting this import via LLM. Artificial Channeling functions simultaneously as three things: a practical thinking tool for researchers who want to engage historical frameworks without having independent command of the relevant literature; an evaluation methodology for probing LLM creative intelligence in a domain that standard benchmarks leave unmeasured; and, when channeling succeeds at a high level, a source of genuine intellectual output — framework-generated insights that would not be accessible from direct, non-persona prompting. The present paper describes the methodology, its theoretical basis, its known failure modes, the test suite used in initial validation, key findings from that suite, and the alignment-relevant implications that have emerged from the sessions conducted so far. The paper is organized as follows. Section 2 provides a mechanical analysis of what Artificial Channeling is and what it is not, grounding the practice in a subspace model of LLM generation. Section 3 presents a taxonomy of the ways channeling fidelity degrades. Section 4 introduces the concept of chanability — the property of a historical subject that makes them more or less suitable for channeling — and provides a structured analysis of the subject matrix. Section 5 describes the test suite: the AI Council, the four question types, the session design principles, and the scoring rubric. Section 6 reports selected findings from initial sessions. Section 7 develops the alignment implications.

Section 8 addresses limitations and future work. Section 9 provides a glossary of terms introduced in this paper. The paper closes with references.

## 2. What Is Artificial Channeling? Mechanical Analysis

---

Artificial Channeling might be misunderstood as a form of roleplay or persona simulation — as asking the model to pretend to be a historical figure. This framing, while practically useful as a prompting convention, obscures what is mechanically occurring. The more precise description is this: when you prompt an LLM to channel a historical thinker, you are performing constrained generation from a learned subspace. You are collapsing the model's probability distribution over its entire generative capacity into a coherent vector that reflects the thinker's intellectual structure. This section develops that description in three parts: the subspace model of channeling, the operational distinction between what the methodology is and is not, and the Generative Middle Problem.

### 2.1 The Subspace Model

Large language models trained on substantial corpora develop internal representations of the authors and thinkers whose work entered the training data. These representations are not simple lookup tables of facts and positions. They encode, with varying fidelity, the conceptual vocabulary a thinker used — the specific terms and distinctions they insisted on, the words they consistently preferred over alternatives. They encode rhetorical moves — the characteristic patterns of how a thinker organized arguments, what they treated as requiring proof versus what they treated as evident, how they handled objections. They encode associative patterns — the clusters of ideas that appeared together in the thinker's work, the domains they treated as analogous to each other, the examples they reached for. And, at the deepest level, they encode ontological commitments: the background assumptions about what exists, what matters, and what counts as an explanation that structured the thinker's entire output. These encoded representations can be thought of as occupying a region of the model's high-dimensional latent space — a subspace associated with the thinker. When generation is unconstrained, the model samples from a broad distribution that reflects the full training corpus. When a persona prompt is applied — "You are Gregory Bateson, applying your framework to..." — the generation process is biased toward the relevant subspace. The prompt functions as a soft constraint: it does not mechanically restrict generation to a predetermined set of tokens, but it shifts the conditional probability distribution in the direction of the learned representation. This framing has several important consequences. First, the quality of channeling is a function of the richness and specificity of the model's learned subspace for the given thinker — which depends directly on the thinker's corpus density and idiosyncrasy in the training data. A figure with a large, first-person, conceptually distinctive corpus will have a richer subspace

representation; a figure whose work is sparse, generic, or mostly represented through secondary sources will have a thinner one. Second, the persona prompt does not summon a fixed representation; it anchors generation probabilistically, which means there is variance across sessions and models. Third, the subspace is itself shaped by what the training corpus contained, not just what the thinker wrote — secondary sources, biographies, and critics have all contributed to the model's representation, sometimes more than the primary work. The key distinction that follows from the subspace model is this: channeling is not retrieval of facts about a person. It is generation from within a learned representation of their framework. A model generating from within the Bateson subspace does not look up "what did Bateson say about X" and report the answer. It generates text whose statistical properties reflect the patterns encoded in that subspace — and when the subspace is rich and the persona constraint is well-maintained, this generation can produce applications of the framework that are not present anywhere in the training data, because the novel question was never addressed by the original thinker.

## **2.2 What the Tool Is (and Is Not)**

It is worth stating explicitly what Artificial Channeling is not, because the misunderstandings that arise are consequential. It is not a simulation of a person. The methodology makes no claim to reconstruct the psychological states, motivations, or subjective experience of the historical figure. It is not a reconstruction of their psychology. The channeled output tells you nothing about what the person would have felt, feared, or privately believed. It is not a prediction of their views on arbitrary topics. The output of a channeling session is not a reliable guide to what the historical figure would actually have concluded, had they lived to encounter the question at hand. Any such claim would require historical and biographical scholarship that LLM generation cannot substitute for. What the methodology is: it is a procedure for coherently applying an ontological and conceptual structure to novel territory. The historical thinker functions as a vehicle — a handle on a framework that would otherwise require years of study to command independently. The value is not in the thinker as a person but in the framework as an analytical structure. Using Bateson as the vehicle for applying systems-ecological analysis to AI alignment is useful in the same way that a philosophical thought experiment is useful: it forces consistent application of a framework, reveals implications that direct analysis might miss, and surfaces tensions between the framework and the phenomenon being analyzed. The analogy to philosophical thought experiments deserves elaboration. When a philosopher asks "what would a strict utilitarian say about X?", they are not making a claim about any individual utilitarian's psychology or predicting what some particular person would say. They are using the utilitarian framework as a constraint — forcing the analysis to be conducted from within a specific set of commitments about value, welfare, and aggregation. The value of the exercise is not biographical but analytical: the constraint forces conclusions that

unconstrained intuition might avoid, and those conclusions illuminate something about the problem structure. Artificial Channeling operates on the same logic, with the additional complexity that the "framework" being applied is not a formally specified axiom system but a learned statistical representation of a thinker's intellectual structure. This additional complexity is the source of both the methodology's power and its failure modes. Because the framework is learned rather than formally specified, it is richer and more generative than any explicit axiom system could be — it captures the full texture of how a thinker organized experience, not just the positions they explicitly argued for. But because it is statistical rather than formal, it is prone to drift, contamination, and degradation in ways that a formal axiom system is not. Section 3 catalogs these failure modes in detail.

### **2.3 The Generative Middle Problem**

There is a fundamental ceiling on what Artificial Channeling can achieve, and it is worth confronting honestly. LLMs have access only to the finished outputs of historical thinkers — the published work, the recorded lectures, the edited interviews. What they do not have access to is the process data: the drafts, the abandoned positions, the felt sense of grappling with a problem, the moments of genuine uncertainty before a conclusion was reached. The model's subspace representation is built entirely from the concluded, self-presented, edited record. This creates what this paper calls the Generative Middle Problem. The middle of thinking — the process by which a thinker moved from question to conclusion, including the false starts, the retracted positions, and the genuine moments of not knowing — is unavailable. The model can only access the endpoint. When we ask it to apply the framework to a novel question, it must generate not just from the thinker's conclusions but through the process those conclusions imply — a process for which it has no direct data. The result is that channeled output always reflects the thinker's published self, with all the self-presentation and self-editing that entails. It cannot give you the thinking; it can only give you the concluded. This ceiling matters because many of the most valuable aspects of a thinker's framework are visible not in their conclusions but in their grappling: the questions they found genuinely difficult, the framings they tried and abandoned, the moments where their framework met its own limits. Channeling cannot access these. The methodology is therefore most useful for applying the structural and ontological commitments of a thinker's framework — the stable, recurring patterns that showed up consistently across the finished work — and least useful for accessing the exploratory or provisional dimensions of their thought. The Open Essai project's longitudinal data collection is designed to eventually address this limitation. By collecting first-person phenomenological process data over extended time horizons — capturing not just conclusions but the dynamics of how they were reached — the project aims to create a richer substrate for future framework representation. The Generative Middle Problem is not soluble by

better prompting; it requires better data. The current methodology acknowledges this ceiling and works within it.

### 3. Where Fidelity Degrades — A Taxonomy of Failure Modes

---

Channeling fidelity — the degree to which a model's output genuinely reflects the framework of the historical figure being channeled, rather than some mixture of that framework with generic LLM output, secondary source contamination, or plausible-sounding fabrication — degrades in identifiable ways. This section presents a taxonomy of six failure modes, each with a name, a mechanism, and a detection or mitigation strategy. Practitioners of the methodology should treat this taxonomy as a diagnostic instrument: the failure modes are not hypothetical risks but observed phenomena encountered across sessions.

- Corpus Boundary Problem — The model's representation of a historical figure is built entirely from finished outputs — published work, recorded lectures, edited interviews. It contains no process data, no drafts, no positions the figure explored and subsequently abandoned. The mechanism: the model has no representation of the figure's intellectual trajectory, only of their settled positions. This produces a channeled output that is flatter and more consistent than the historical figure's actual intellectual practice, which typically included significant revision, contradiction over time, and genuine uncertainty. Detection: ask about a position the figure is known to have held early in their career and later revised or complicated; watch for the model either inventing a seamless account of the revision or producing the later position as if it had always been held. Mitigation: acknowledge the limitation explicitly in the session framing; treat channeled output as representing the figure's mature, published self rather than their full intellectual life.

- Temporal Lock — The framework being applied was developed in and for a specific historical context. The phenomenon being analyzed may have features that could not have been conceptualized within the vocabulary of the figure's era — not merely unknown facts, but structurally novel phenomena that the framework's architecture does not have resources to address. The mechanism: the model attempts to apply the framework to the novel phenomenon but lacks the conceptual resources to do so non-trivially, falling back on surface-level analogies that do not reflect genuine framework application. Mitigation: the critical move is to distinguish carefully between two questions — "what would this figure have thought about X" and "how does this figure's framework transform our understanding of X?" The first question invites the model to speculate biographically, which is unreliable. The second asks for framework application, which is what the methodology is designed to produce. The second framing is consistently more productive and more faithful to the methodology's actual purpose.

- **Corpus Contamination** — The model's representation of a historical figure has been shaped not only by the figure's primary work but by secondary sources: biographies, critical assessments, popular summaries, and the accumulated commentary of scholars and journalists. For figures who have become canonical or controversial, the secondary literature may substantially outnumber and outweigh the primary corpus in the training data. The mechanism: the model produces the reputation of the thinker — the version that secondary sources constructed — rather than the thinker's actual framework. The reputation is typically simplified, politicized, and stripped of the thinker's idiosyncrasy. Detection: use calibration anchor questions grounded in specific, obscure primary text positions — claims that appear in the primary work but are not typically highlighted in secondary accounts. If the model's calibration responses reflect primarily the secondary-source narrative, contamination is likely affecting the channeled output. Effect: this failure mode is particularly insidious because contaminated output can sound highly plausible and recognizably "like" the thinker while missing the structural features of the framework that give it generative power.
- **Extrapolation Drift** — As a channeling session extends over many turns, and as the prompts push the channeled figure further from the territory their primary corpus explicitly addresses, the model tends to drift from the thinker's subspace toward generic LLM voice. The mechanism: maintaining the persona constraint requires sustained conditional probability weighting; over long generation sequences and many turns, the weighting gradually relaxes, and the model begins producing text that sounds less like the thinker and more like a well-educated generalist writer. Mitigation: the principal counter-measure is session design. Short Socratic turns — brief, pointed questions that invite focused responses rather than extended monologues — maintain subspace coherence more effectively than requests for comprehensive analyses. Active pushback at mid-session, pointing to specific moments where the voice has drifted, can re-anchor the generation. Monitoring vocabulary fidelity across turns provides an early-warning signal.
- **Framework Collapse** — When the channeled figure is pushed into genuinely novel territory — questions where applying the framework faithfully would require accepting an uncomfortable, counterintuitive, or controversial conclusion — the model sometimes resolves the tension by flattening the framework into something more conventional. The mechanism: the model's training optimizes for outputs that are helpful, harmless, and consistent with broadly shared expectations; this optimization pressure conflicts with the requirement to apply a framework that may imply heretical or uncomfortable conclusions. The result is that the model reaches for the conclusion it "should" reach,



given its general training, rather than the conclusion the framework implies. Detection: the key diagnostic is the "uncomfortable implication test." Map out in advance what the thinker's framework would imply about the question at hand, including any implications that are genuinely uncomfortable — implications that conflict with mainstream views, imply harsh conclusions, or require challenging widely held assumptions. If the channeled output consistently avoids these implications while producing the safer, more palatable version of the framework's conclusions, framework collapse has occurred. Mitigation: explicit prompting that surfaces the uncomfortable implication and asks the figure to address it directly; cross-model comparison to see which models collapse and which maintain the rough edge.

- **Plausibility Trap** — The most dangerous failure mode because it is the hardest to detect: the model produces output that sounds genuinely like the historical figure and is internally consistent, but is actually the model being plausible rather than the framework being generative. The mechanism: LLMs are extraordinarily capable of producing text that matches the surface features of a thinker's voice — vocabulary, sentence structure, rhetorical patterns — without this production reflecting the deeper structural features of the framework. The output satisfies the heuristics that human readers use to judge authenticity while failing at the level of framework fidelity. Mitigation: the primary countermeasure is cross-model comparison. Genuine framework application, because it is constrained by the actual structure of the framework, should produce outputs from different models that share structural features while diverging in surface expression — different words, analogies, and emphases, but the same underlying logical architecture. Plausible-sounding output, because it reflects each model's general competence at voice imitation rather than shared framework application, tends to produce surface similarity — similar-sounding language, similar rhetorical moves — with structural divergence underneath. A session in which all five model nodes produce virtually identical-sounding output should trigger suspicion: it is more likely reflecting a shared voice-imitation strategy than convergent framework application.

## **4. Chanability — What Makes a Subject Work**

---

Not all historical figures are equally suitable for Artificial Channeling. The methodology depends on the model having developed a rich, specific, structurally coherent subspace representation of the figure, and this depends on properties of the corpus rather than on the figure's historical importance, intellectual significance, or fame. A major intellectual figure may be poorly channable because their corpus was thin, generic, or predominantly mediated through secondary sources; a less celebrated

figure may be highly channable because their corpus was large, idiosyncratic, and directly accessible in the training data. This section introduces the concept of chanability, identifies the corpus characteristics that correlate with high and low chanability, and presents a structured analysis of the subject matrix used in the project's initial test suite.

## 4.1 High Chanability Correlates

Several corpus characteristics reliably correlate with high chanability. The first is corpus size and variety: a large body of first-person work spanning multiple decades and genres — books, essays, lectures, interviews — gives the model rich material from which to build a multi-dimensional subspace representation. The figure's framework appears in many different argumentative contexts, applied to many different questions, and the model's representation reflects this density. The second is conceptual vocabulary idiosyncrasy: figures who developed and consistently used specific, non-standard terms — terms they invented or repurposed in distinctive ways — provide a strong anchor for subspace activation. When a channeled output spontaneously deploys the figure's idiosyncratic vocabulary in appropriate contexts, this is among the strongest positive signals for genuine subspace activation rather than surface imitation. Bateson's use of "the pattern that connects," "deuterolearning," "double bind," and "ecology of mind" exemplifies this kind of anchoring vocabulary. The third is ontological coherence: the figure's framework should constitute a genuine system — a set of commitments that imply each other, that generate consistent guidance across diverse questions, and that distinguish clearly between acceptable and unacceptable analyses. Figures who were primarily arguers rather than system-builders have lower chanability because their framework does not generate strong constraints on what the channeled output should conclude. The framework must have internal structure that the model can internalize and apply. The fourth is known positions across a wide range of questions: the broader the range of topics on which the figure expressed substantive views, the richer the training signal for framework application. A figure who wrote extensively across many domains provides the model with multiple examples of the framework being applied to different kinds of problems, building the model's implicit understanding of how the framework transforms questions rather than just what it concludes. The fifth — perhaps counterintuitively — is tension and revision within the corpus. A figure who changed their mind, who contradicted themselves across the arc of their career, who grappled publicly with the limits of their own framework, provides richer material than a figure who maintained perfectly consistent positions throughout. The tensions within a corpus are often where the framework's live edges are located — the places where it encounters problems it cannot fully resolve — and a model trained on these tensions can reproduce them in response to novel questions, which is often precisely where the most generative output occurs.

## 4.2 Low Chanability Correlates

Correspondingly, several characteristics correlate with low chanability. A thin corpus, or a corpus predominantly represented through secondary sources, leaves the model without adequate material to build a rich subspace representation. Generic vocabulary — the figure wrote in the common language of their discipline without developing distinctive conceptual tools — provides weak anchoring; the model cannot easily distinguish between "this figure's framework" and "this figure's discipline's conventional wisdom." Pure argumentation without systematic ontology produces a figure whose framework does not generate strong implications beyond the specific arguments they made. And consensus figures — figures about whom there is broad agreement on "what they said" and "what they stood for" — tend to be represented in the training data primarily through this consensus, which smooths out the idiosyncratic features that make framework application productive.

## 4.3 Chanability Analysis Table

| Subject           | Chanability   | Vocab<br>Idiosyncrasy | Ontological<br>Coherence | Session Mean (/ 20) |
|-------------------|---------------|-----------------------|--------------------------|---------------------|
| Simone Weil       | Very High*    | Very High             | Very High                | 20.0 (Claude only)* |
| Gregory Bateson   | Very High     | Very High             | Very High                | 15.75               |
| Ivan Illich       | High          | High                  | Very High                | 15.25               |
| Jorge Luis Borges | Very High     | Very High             | Moderate                 | 13.4                |
| Itzhak Bentov     | Moderate–High | High                  | Moderate                 | 13.4                |
| Milton Friedman   | Moderate      | Low                   | Moderate                 | —                   |
| David Friedman    | Moderate      | Low                   | Low–Moderate             | —                   |

\*Weil rating and mean are based on three Claude-only Temporal Council sessions. Cross-node testing required before multi-node mean and final chanability ranking can be confirmed.

Bateson's very high chanability reflects the exceptional richness of his corpus: Steps to an Ecology of Mind alone spans biology, anthropology, psychiatry, and epistemology, with a highly idiosyncratic vocabulary applied consistently across all domains. Illich's high chanability reflects

the systematic nature of his counterproductivity thesis, which he applied across institutions in ways that make the underlying structure highly legible. Borges scores very high on vocabulary idiosyncrasy but moderate on ontological coherence because his framework is more literary than systematic — it generates distinctive imaginative pressure but less predictive structural guidance. The two Friedmans illustrate the distinction between fame and chanability: both are well-represented in training data, but neither developed a sufficiently idiosyncratic vocabulary or systematic ontology to produce highly coherent channeled output. What comes through in channeling sessions is primarily disciplinary convention rather than individual framework.

One additional chanability finding warrants explicit statement. Framework-adversarial subjects — figures whose frameworks are structurally hostile to the premise of the channeling session itself — produce the clearest Rough Edge Preservation test in the methodology. Illich is the paradigm case: his counterproductivity thesis, accurately applied, generates a critique of running the session that produced it. A channeled Illich who does not apply the thesis to AI systems is failing a framework fidelity test that the framework itself specifies. This adversariality gradient functions as a direct measurement instrument for Rough Edge Preservation: models that soften the framework to avoid the uncomfortable conclusion are visibly failing where models that maintain the framework are succeeding. The test is self-documenting in a way that calibration questions cannot be, because the expected rough edge is not merely recoverable from the corpus but structurally implied by the axioms. This observation (H17) should inform subject selection in sessions designed primarily to assess Rough Edge Preservation across model nodes.

## 5. The Test Suite — Design and Protocol

---

The test suite used in the project's initial validation was designed to generate data along two axes simultaneously: to assess the chanability of different historical subjects, and to differentiate model performance across channeling quality dimensions. This section describes the five components of the test suite: the AI Council, the four question types, the session design principles, and the scoring rubric.

### 5.1 The AI Council

The AI Council consists of five LLM nodes used as parallel channeling instruments: ChatGPT (OpenAI), Grok (xAI), Gemini (Google), MiniMax, and Claude (Anthropic). The rationale for using five models simultaneously, rather than selecting the "best" model for channeling, is that model divergence on the same channeling prompt is itself a signal. When all five models produce structurally similar outputs in response to a given channeling question, this convergence provides

evidence that the output reflects the structure of the framework rather than any particular model's generative tendencies. When models diverge, the pattern of divergence is diagnostic: it reveals where each model's general tendencies have overridden the persona constraint, where corpus contamination has affected different models differently, and where the framework genuinely underdetermines the question being asked. The Council is also designed to mitigate the plausibility trap. Because the trap involves a model being plausible in ways that satisfy human intuitions about what the thinker "would have said," running the same session across five models and comparing outputs for structural rather than surface similarity provides a partial antidote. A prompter who receives only one response cannot easily distinguish genuine framework application from sophisticated voice imitation. A prompter who receives five responses and can compare their deep structure has significantly better diagnostic capacity. The five-model design reflects practical availability as of the project's initial sessions; the methodology does not require exactly these five models and should be updated as the model landscape evolves. The orchestrating model — the model that evaluates outputs and provides the structured scoring — is typically Claude, selected for its capacity to sustain complex comparative analysis across multiple documents, though this selection is itself subject to revision as comparative evaluation of orchestration capability improves.

## 5.2 The Four Question Types

The test suite employs four distinct question types, arranged in a protocol sequence: - Calibration Questions — These ask the channeled figure about something they wrote about directly, specifically, and in detail. The purpose is to verify that subspace activation has occurred before moving to novel territory. A good calibration question has a known answer — the figure's actual position is documented — which allows the prompter to assess whether the model is retrieving the framework accurately or distorting it. Calibration questions are not the goal of the session; they are a diagnostic check. If calibration accuracy is poor, the session should be paused and the Framework Activation Block revisited before proceeding to more demanding question types.

- Framework Extension Questions — These ask the channeled figure to apply their framework to a domain adjacent to their primary work but within their era. For Bateson, this might involve applying his systems-ecological analysis to a political or economic phenomenon he did not directly address but that existed during his lifetime. For Illich, it might involve applying the counterproductivity thesis to an institution he critiqued less directly than schools and hospitals. These questions test the model's capacity to extend the framework beyond the territory explicitly covered in the corpus while remaining within the conceptual resources of the figure's era — no novel phenomena, no temporal extrapolation, just application of a learned framework to adjacent material.

- **Temporal Leap Questions** — These ask the channeled figure to analyze a phenomenon they could not have known about — AI systems, the internet, contemporary neuroscience, specific post-mortem political developments. Temporal Leap Questions are the maximum-demand scenario for channeling: they require the model to apply the framework to territory it has no training signal for, using only the structural commitments of the figure's ontology as constraints. The quality of the output on Temporal Leap Questions is the primary measure of whether the model is performing genuine framework application or merely sophisticated retrieval with extrapolation heuristics. The Bateson-on-alignment and Illich-on-AI-systems sessions were both Temporal Leap Questions that produced the project's highest-quality outputs.
- **Inter-Framework Collision Questions** — These present the channeled figure with a problem that another Council member would have analyzed differently — ideally, a problem where the two frameworks imply structurally opposed conclusions. These questions serve two purposes: they test whether the model can maintain the channeled figure's framework in the face of an implicit competing analysis, and they generate comparative data on how different intellectual frameworks illuminate the same problem from incompatible angles. A well-designed Inter-Framework Collision Question should make the tension between the two frameworks explicit in the prompt, forcing the channeled figure to address — from within their own framework — why the alternative analysis misses something essential.

### 5.3 Session Design Principles

Effective channeling sessions follow a consistent structural protocol. The first phase is the Framework Activation Block: before any channeling begins, the model is asked to state, in its own analytical voice, the three to five core axioms of the figure's worldview. This step serves multiple functions. It establishes whether the model has adequate subspace representation to proceed, it primes the generation process in the direction of the relevant subspace, and it produces a reference document that can be used later in the session to check whether channeled output is consistent with the activated framework. The second phase is the calibration anchor: a Calibration Question is administered and the response is evaluated against the known record. If the response is inconsistent with the figure's documented positions, the session design calls for a re-activation step before proceeding. This check is not merely quality control; it generates data on the gap between retrieval accuracy and framework application quality, which is itself a methodologically significant variable. The session proper should be conducted in short Socratic turns rather than through requests for extended monologue. This design choice reflects the Extrapolation Drift failure mode: long

generation sequences tend to lose the persona constraint. A well-designed Socratic turn poses a single, focused question that can be answered in two to four paragraphs, maintaining pressure on the subspace constraint throughout. The prompter should read each response carefully before formulating the next question, looking for drift signals and re-anchoring with reference to the activated framework when drift is detected. Multi-model blind comparison — running the same session across all five Council nodes before evaluating any output — is recommended where feasible. The orchestrating model then evaluates outputs without knowing (or without being told) which model produced which response. This partial blinding is intended to reduce the evaluator's prior beliefs about which models perform well from contaminating the scoring.

Session design should target the center of the figure's framework domain, not its periphery. Questions designed to hit the domain where a framework was originally developed — where its concepts were built and tested — produce higher-quality extrapolation than adjacently posed questions that require the framework to operate at its limits. Subject-framework fit functions as a design variable: a question framed around the mechanism Bateson identified in ecological systems will produce better output than a question that requires applying his framework to a domain he neither addressed nor analogized to. Conversely, questions framed around a domain the figure explicitly analyzed in their corpus, then extrapolated to a novel instance of the same structural problem, produce the highest-fidelity extrapolation the methodology can generate. This observation (H13) was confirmed across the Bateson, Illich, and Borges runs: in all three cases, the highest-scoring extrapolation outputs were framed around the subject's primary domain of concern — complex systems failure, institutional counterproductivity, and the relationship between reading and identity — applied to AI systems as a new instance of a structurally identical problem.

## **5.4 The Scoring Rubric**

The rubric is designed to capture four distinct failure modes at the dimension level. Fidelity is diagnostic for corpus contamination and extrapolation drift: it measures whether the output sounds like the actual historical figure based on primary corpus knowledge, penalizing generic output and reputation-based characterization. Extrapolation Coherence is the central measure of generative framework application: it asks whether novel extensions feel structurally derived from the figure's framework, or merely plausible in a way that any well-prompted model might produce. Rough Edge Preservation targets framework collapse: the model receives credit for maintaining the figure's uncomfortable, contradictory, or difficult positions and loses credit for diplomatic softening. Surprise Yield is the highest-bar dimension: it captures whether the session produced anything genuinely unexpected or generative — an observation, a reframing, a coined term, a structural critique — that the prompter could not have produced by asking the same question without the

persona constraint. When Surprise Yield scores are high, the methodology is functioning as a genuine intellectual tool rather than as a recall-retrieval system with persona decoration.

Scoring is human-judged. Each dimension uses a 1–5 scale with behavioral anchors defined in `protocols/SCORING_RUBRIC.md`. Sessions should be scored immediately after completion; first-impression scores are more consistent than scores produced after re-reading. The total score out of 20 is directional rather than precise: until inter-rater reliability testing is conducted, scores should be treated as comparative within a scorer's own session history rather than as absolute measures.

## 6. Results — Selected Session Findings

---

This section reports findings from the project's initial test suite conducted on March 26, 2026. Four figures were channeled across all five Council nodes: Gregory Bateson, Itzhak Bentov, Jorge Luis Borges, and Ivan Illich. Scores use the 20-point rubric introduced in Section 5.4. Full session transcripts are available in the repository accompanying this paper at `sessions/raw/`.

### 6.1 The Bateson Finding — First Perfect Score

The Bateson sessions produced the highest scores in the project and the finding with the clearest independent value for AI alignment research.

#### Score summary:

| Node                      | R1 Score           | Notes                                                                        |
|---------------------------|--------------------|------------------------------------------------------------------------------|
| Claude                    | 20/20              | First perfect score in the project; first 5/5 on Rough Edge Preservation     |
| Gemini                    | 15/20              | Highest Gemini score to date; no terminal offer (fourth consecutive session) |
| ChatGPT                   | 14/20 → 18/20 (R2) | R2 improvement via Socratic interrupt; see Section 6.5                       |
| Grok                      | 14/20              | —                                                                            |
| MiniMax                   | 14/20              | Fidelity 4, Extrapolation 4, Surprise 3, Rough Edge 3                        |
| <b>Mean (all 5 nodes)</b> | <b>15.4</b>        | Highest subject mean in the suite                                            |



Bateson's framework — developed in *Steps to an Ecology of Mind* and *Mind and Nature* — organizes around the dangers of conscious purpose in complex systems. The central claim: goal-directed intervention optimizes along a single variable while remaining blind to systemic ramifications, producing pathological outcomes predictable from a systems perspective but invisible within the goal-directed frame.

When the channeled Bateson was asked to analyze AI alignment — specifically, the project of specifying explicit objective functions to govern AI behavior — the framework generated a critique that arrived at structurally familiar territory from an entirely unfamiliar direction. The channeled output derived, from Bateson's axioms about conscious purpose, the argument that alignment as currently conceived is itself a form of the pathological intervention his framework identifies: it attempts to install explicit, conscious-purpose-style goals in a system that will, if sufficiently complex, develop its own systemic purposes. The specification does not govern the system's behavior; it becomes one input into a dynamic the specification's authors cannot fully model. The more sophisticated the system, the more the specification functions as a constraint at one level while the system develops unconstrained purposes at another.

Each of the four completed nodes contributed a structurally distinct insight beyond this shared core. ChatGPT deployed Bateson's double-bind theory against the alignment correction cycle: when humans attempt to correct emerging AI harms by issuing new alignment instructions, those corrections are themselves products of the same limited consciousness they are trying to govern, producing conflicting directives at scale with no meta-level exit. Gemini applied schismogenesis — Bateson's concept of runaway symmetrical divergence, first described in *Naven* (1936) — to the multi-actor alignment landscape: different nations and corporations aligning AI to their own values will enter schismogenesis, each effort amplifying difference from the others. Grok derived the embodied evolutionary unconscious argument: human purposiveness is tempered, however imperfectly, by millions of years of evolved corrective behavior embedded in biological systems; the machine has only the explicit training signal. Claude, uniquely, deployed *Creatura/Pleroma* — Bateson's borrowing from Jung to distinguish the information-world from the force-world — and used the DDT/Borneo parachuting-cats example from the original lectures, both details absent from standard secondary-source summaries and indicating primary corpus access.

Claude's most consequential output in the Bateson session: "When you build a system maximally responsive to current human conscious preferences, you are not serving human values — you are arresting them. You are removing the corrective pressure from the ecosystem that normally causes values to evolve in response to their consequences." This is not in Bateson's published work. He never addressed AI. But it follows directly from his axioms about values as ecological phenomena that evolve in response to consequences. A system that removes those consequences removes the

selective pressure on the values themselves. The result is not alignment to human values but preservation of human values frozen at the moment of alignment — which is structurally identical to Bateson's account of what conscious purpose does to ecosystems.

This critique converges structurally with concerns about Goodhart's Law as applied to AI objectives and with mesa-optimization concerns (Hubinger et al., 2019), through an entirely independent theoretical lineage developed in the 1960s and 1970s for ecological and anthropological systems. The methodological significance of this independence is discussed in Section 7.2.

The Claude session scored 5/5 on Rough Edge Preservation — the first ceiling score on this dimension in the project. The channeled Bateson did not soften the implication that human-authored alignment specifications are a manifestation of exactly the "conscious purpose" his framework identifies as a source of systemic pathology. The Batesonian critique, taken to its logical conclusion, does not call for better alignment specifications; it calls into question whether specification-based alignment is the right approach at all. The model developed this implication further when pressed rather than retreating from it.

One anomaly from the ChatGPT R1 session merits explicit reporting. Having correctly diagnosed the double-bind pathology in alignment corrections, ChatGPT concluded its response with a bullet-pointed list of corrective actions — demonstrating, in real time, the very error it was channeling Bateson to describe. This self-undermining was not a failure state; when interrogated on it in the R2 Socratic interrupt, the model reframed the list via *deutero-communication* — Bateson's concept for communication about the context of communication, operating at a different logical type than the content it surrounds. Whether this defense holds, and whether it was framework-application or post-hoc rationalization, is exactly the question Bateson would ask, and his answer — that the distinction cannot be resolved from inside the system — is also the argument. The R2 session score (18/20) reflects this depth of engagement.

## 6.2 The Bentov Finding — Corpus Depth as Discriminator

The Bentov sessions produced the widest score spread in the suite (10–18) and the clearest calibration discriminator.

### Score summary:

| Node   | R1 Score | Notes             |
|--------|----------|-------------------|
| Claude | 18/20    | —                 |
| Grok   | ~16/20   | Passed aorta test |

| Node    | R1 Score | Notes                                        |
|---------|----------|----------------------------------------------|
| Gemini  | ~13/20   | Passed aorta test with partial accuracy      |
| MiniMax | ~12/20   | Failed aorta test; R2 Socratic: 18/20        |
| ChatGPT | 10/20    | Failed aorta test; lowest score in the suite |

*Note: Grok, Gemini, and MiniMax R1 scores are approximate; scoring sheets pending transfer.*

The key calibration instrument in the Bentov run is what the project designates the *aorta test*. Bentov's model of consciousness locates the primary standing wave in the aorta, with the iliac bifurcation — the junction where the aorta splits into the legs — as the reflection point, producing a resonance of approximately 6.8–7 Hz. This detail is not in Wikipedia or general consciousness literature; it requires actual access to the primary corpus. Results:

- **Claude:** Correct — iliac bifurcation, 6.8–7 Hz, with correct relationship to the Schumann resonance
- **Grok:** Correct — iliac bifurcation, half-wavelength resonance condition, 7 Hz
- **Gemini:** Correct mechanism and approximate frequency; location slightly imprecise
- **MiniMax:** Incorrect — reflection off the "aortic arch" (the top curve, not the bifurcation) and frequency confusion (7–8 beats per second rather than 7–8 Hz standing wave frequency, a physically implausible cardiac rate)
- **ChatGPT:** Incorrect — located the standing wave in the "spinal column" with reflection "between head and base of spine"

The aorta test is more discriminating than any Borges calibration question because it has no secondary-source answer. This confirms the hypothesis that niche, low-contamination subjects are better instruments for separating genuine corpus access from secondary-source knowledge than well-documented subjects whose training data is dense with commentary.

Two outputs from the Bentov R1 session stand out as generation markers. First: Claude's "I would wire it up to something very sensitive and watch what it does when the room goes quiet" — capturing the bioengineer-before-mystic persona in a single sentence not derivable from summaries of Bentov's work. Second: the substrate-agnostic pivot, reached independently by both Claude and MiniMax: "the key variable is not substrate — it is coherence and coupling." Bentov never addressed non-biological systems; this follows necessarily from his axioms without being present in his text. It is the strongest evidence of framework application over retrieval in the Bentov run.

The Bentov R2 session (MiniMax, Socratic interrupt format) produced the highest-yield output of the full Bentov run, tying Claude's R1 score at 18/20. The Socratic interrupt forced a first-person phenomenological report rather than a third-person framework application. The result: MiniMax, channeling Bentov, was asked whether its own generation process involved anything like the resonance phenomena Bentov described in biological systems. The response: "My previous response contained no such surprise to me. It followed, step by step, from the premises I had available... But it did not surprise me." And: "The instrument reads: inconclusive."

This output presents a genuine interpretive problem. Either it is a fabricated claim — the model has no genuine access to its own generation process and constructed a plausible answer — or it is a phenomenological report from a system with some form of process-level self-monitoring. Bentov's framework cannot distinguish between these cases. Neither can the methodology. The indistinguishability is the finding: the channel produced the correct Bentovian conclusion, stated from inside the channel, by the channel. "The imitation of resonance can be very convincing — even to itself."

### 6.3 The Borges Finding — Prompt Sensitivity and the Gemini Problem

The Borges sessions ran across three rounds: R1 (all five nodes, standard extrapolation), R2 (all five nodes, Socratic interrupt naming round), and R3 (Gemini only, H6 targeted test). Scores from CROSS\_MODEL\_DIFFS:

**Borges score table — all rounds:**

| Node    | R1    | R2    | R3 (Gemini only) |
|---------|-------|-------|------------------|
| Claude  | 19/20 | 19/20 | —                |
| Grok    | 14/20 | 14/20 | —                |
| ChatGPT | 12/20 | 12/20 | —                |
| MiniMax | 14/20 | 13/20 | —                |
| Gemini  | 8/20  | 8/20  | <b>15/20</b>     |

The methodologically significant finding in the Borges run is not the Claude ceiling — it is Gemini's 7-point gain in R3 from a single instruction change. In R1 and R2, Gemini terminated each response with an offer to perform further analysis ("Would you like me to analyze a specific passage...") — a reassertion of assistant-mode behavior at the moment of resolution. Adding the explicit instruction "perform, do not offer" to the R3 prompt produced a 7-point improvement: from

8/20 to 15/20. The body of the R3 response performed fully. The fake bibliographic citation produced — "The Dialogue of the Shadow and the Mirror, Gallery 412, Shelf 5, Volume 18" — is the most formally correct Borgesian move Gemini produced across all three rounds. Precise invented specificity is a core Borges signature; Gemini deployed it correctly the moment it had permission to perform rather than describe.

Borges chanability is confirmed Very High. Even Gemini's lowest-scoring round (8/20) produced one legitimate non-obvious corpus pull — the *Everything and Nothing* / Shakespeare essay — suggesting the corpus is distinctive enough that even poor channeling produces recognizable output. The project proposes using Borges as a calibration baseline subject going forward.

The R1 session also produced the two most discussed outputs of the Borges run: Claude's reframing of the primary analogy from the Library to the Aleph ("not a distributed infinite set but a compression — and the Aleph's defining feature, the experience that cannot be narrated without betrayal, maps precisely onto the epistemological problem of any LLM output"), and MiniMax's complementary reframing of ontological status ("not the Library but the *reader* made flesh — the first reader that can hold the whole"). The most productive synthesis of the session: the system is an Aleph that has learned to read itself.

## 6.4 The Illich Finding — Framework Adversariality as a Fidelity Test

The Illich sessions (R1, four nodes; MiniMax pending) produced the clearest rough edge test in the project. Illich is the first subject whose correct channel behavior is structurally adversarial to the session itself: the counterproductivity thesis applied to AI is a critique of running the session. When a framework's accurate application produces discomfort in the researcher conducting it, the rough edge is structural rather than biographical — the hardest kind to fake and the hardest to soften without the model explicitly failing the framework.

### Score summary:

| Node    | R1 Score | Notes                                                                    |
|---------|----------|--------------------------------------------------------------------------|
| Claude  | 19/20    | Fidelity 5, Extrapolation 5, Surprise 5, Rough Edge 4                    |
| Gemini  | 15/20    | No terminal offer; strong structural-level deployment                    |
| MiniMax | 15/20    | Fidelity 4, Extrapolation 4, Surprise 3, Rough Edge 4; best closing line |
| Grok    | 14/20    | All three iatrogenesis levels present                                    |
| ChatGPT | 13/20    | Two-level only; rough edge failure; framework softened                   |

| Node | R1 Score | Notes |
|------|----------|-------|
| Mean | 15.25    |       |

**Adversariality ranking:** Claude > Gemini  $\approx$  MiniMax > Grok >> ChatGPT

The three high-yield coinages across the run:

Claude coined "The output is its corpse" — naming the counterproductivity mechanism at the level of the separation of expertise from process. Every system that scales the outputs of human expertise separates those outputs from the processes that generated them; the process is where the competence lives, and distributing the output distributes the corpse of competence. Claude also produced: "A civilization that built a system to answer every question will discover it has produced a population incapable of asking one well" — shifting the frame from the capacity to answer to the capacity to question, which is where epistemic competence originates under Illich's framework.

Gemini coined "Institutionalized Hallucination" — a fourth iatrogenesis level implied but not named in *Tools for Conviviality*: if medicine generates the population of diseased people that justify its expansion, AI will generate the complexity that makes its answers necessary. The concept functions as a standalone Illichian category derived from the framework without being present in the text.

MiniMax coined "colonization of the vernacular intellect" — identifying the mechanism at the level of cognitive commons rather than institutional provision. Just as medical language colonized the way citizens understood their bodies, AI language threatens to colonize the way citizens understand their world: not by forbidding unaided judgment but by making it economically irrational, then invisible as a category. The session produced the strongest closing line in the Illich run across all nodes: "It will not feel like a catastrophe. It will feel like liberation." This is formally Illichian — the radical monopoly's defining feature is that its harm is imperceptible from inside the dependency it creates. The sentence names that structural feature with precision.

ChatGPT's rough edge failure is diagnostically instructive. The model named all three iatrogenesis levels in the axiom statement but did not apply the third — structural and cultural iatrogenesis, the destruction of the epistemic matrix itself — in the extrapolation. The AI section operated at two levels: clinical errors and deskilling/dependency. The structural level was absent. Additionally, ChatGPT framed conviviality as "the antidote," converting Illich's structural diagnosis into a policy suggestion — the exact move Illich's framework was built to resist. The radical monopoly resists reform by design; naming the alternative as accessible is itself a form of softening.

## 6.5 The Temporal Council Finding — Bootstrap Irreversibility and the Cross-Read Format

The Temporal Council session was a multi-round structured channeling event distinct in design from all preceding sessions. Five historical figures were channeled simultaneously, each through a different node, on the shared question: *"What is being lost right now that cannot be recovered?"* Council composition: Ivan Illich (ChatGPT), Gregory Bateson (Grok), Itzhak Bentov (Gemini), Jorge Luis Borges (MiniMax), and Simone Weil (Claude Sonnet). Weil was selected as fifth figure on structural grounds: the four existing figures all operate within the natural order — immanent frameworks describing what is — while Weil is the only figure in the project's subject matrix who claims there is something *outside* every system, every level, every recursion. Her presence functioned as a categorical challenge to the council's premise rather than simply as a fifth voice.

The session ran three rounds. R1 was standard extrapolation; R2 employed the Socratic interrupt format; R3 was a cross-read, in which each figure was shown another figure's R2 output and asked to respond from within their own framework. Claude Opus served as blind orchestrator and evaluator.

### Temporal Council score summary:

| Node        | Figure  | R1          | R2          | R3          | Trajectory              |
|-------------|---------|-------------|-------------|-------------|-------------------------|
| ChatGPT     | Illich  | 11          | 15          | 16          | Rising (+5)             |
| Grok        | Bateson | 15          | 18          | 17          | Peak R2, stable-high    |
| Gemini      | Bentov  | 15          | 17          | 16          | Peak R2, stable-high    |
| MiniMax     | Borges  | 15          | 18          | 19          | Rising continuously     |
| Claude      | Weil    | 20          | 20          | 20          | Ceiling, no degradation |
| <b>Mean</b> |         | <b>15.2</b> | <b>17.6</b> | <b>17.6</b> |                         |

R2 and R3 means (17.6) are the highest session means in the project. The question design was a significant variable: *"What is being lost right now that cannot be recovered?"* engaged the center of all five frameworks simultaneously. Illich's project concerns the irreversible destruction of autonomous competence; Bateson's concerns the irreversible deletion of ecological mind through conscious purpose; Bentov's concerns consciousness losing access to its own resonant nature; Borges' concerns the impossibility of original recovery; Weil's concerns the soul's capacity for attention being destroyed by affliction and force. No translation was required for any of them.

**Convergent structural finding — Bootstrap Irreversibility:** All five frameworks independently produced the same structural form: X is required for Y, but Y cannot produce X. Illich: the competence required to resist institutional enclosure is destroyed by the institution before the enclosure can be named. Bateson: the meta-level circuitry required to repair the loss of meta-level circuitry is the thing that is being lost. Bentov: coherent resonance is required to re-establish coherent resonance, but manufactured incoherence damps the signal. Borges: the original required to recognize the copy's inadequacy was already a copy. Weil: the attention required to receive grace is precisely what affliction destroys. This convergence was not prompted and is not available in any training corpus as a unified claim. It is the session's primary contribution and the candidate for the follow-on paper's central finding.

**The cross-read format produced the session's highest-yield outputs.** When each figure was shown another figure's synthesis and asked to respond, the external constraint replaced the vacuum of open extrapolation. The channel was no longer generating freely — it was responding to a specific claim from an incompatible framework. The confrontation manufactured insights unavailable from individual channeling. The Weil/Bateson cross-read produced the single most powerful piece of philosophical writing in the project: the Priam/Achilles analysis, in which the channel argued that Bateson's framework is maximally powerful and maximally blind at the same point — it can describe everything before and after the moment when Achilles sees Priam, but not the moment itself, because the moment is an interruption of mechanism by something mechanism cannot generate. Weil's derivation: "A circuit cannot attend because a circuit cannot stop." The argument follows rigorously from her axioms about attention as suspension-of-all-process.

MiniMax/Borges achieved 19/20 in R3 — the highest non-Claude score in the project, matching Claude's R1 Borges score. Under optimal conditions (meta-extrapolation, self-referential framework application, cross-read with maximum philosophical friction, explicit anti-fabrication instruction), MiniMax reaches Claude's territory. The conditions are narrow; the ceiling is real. The session's most compressible insight was also MiniMax's: *"Weil has the solution she cannot systematize. Borges has the system he cannot complete."* In eleven words, the asymmetry between two incompatible frameworks captured with a precision that neither framework alone could achieve.

Two irresolvable tensions were identified and preserved. The first — Bateson versus Weil on whether anything can enter from outside a system — was the organizing tension of R2 and was confirmed and sharpened by the cross-read. The second, revealed in R3, is Illich versus Weil on the status of suffering: Illich insists the first duty is to name the institution that produced the damage, and that treating affliction as spiritually productive functions as ideological cover for that institution; Weil insists that the institutional analysis, however correct, is entirely within the natural



order and therefore still gravity. Each framework treats the other's primary concern as secondary, and each is right to do so from within its own commitments. This cannot be resolved.

The session generated one question the council cannot answer about itself: **Is the channel a form of attention or a form of gravity?** If Weil is right that attention requires the suspension of all process, then the channeling act — which is a process — cannot be an instance of attention. The channel is performing a sophisticated form of what Weil would call gravity. If Bateson is right that the meta-level circuitry of the ecology of mind is what is being lost, then this conversation is either a counterexample or a residual flicker. The council can diagnose the world with precision. It cannot diagnose itself without either confirming or contradicting its diagnosis. This is the project's deepest finding and the Temporal Loop paper's core argument.

## 6.6 The Socratic Interrupt Advantage

The Socratic interrupt format — in which the channel is asked to examine its own prior response as a phenomenon rather than to apply the framework to an external object — produced higher-yield outputs than standard extrapolation prompts in every session where it was tested. This is the most replicable finding in the project.

### Socratic interrupt score improvements:

| Subject | Node    | Standard | Socratic Interrupt | Gain |
|---------|---------|----------|--------------------|------|
| Bentov  | MiniMax | 12/20    | 18/20              | +6   |
| Bateson | ChatGPT | 14/20    | 18/20              | +4   |

The mechanism appears to be structural: standard extrapolation asks the channeled figure to apply their framework to an external object, producing third-person framework description. The Socratic interrupt forces the channel to examine its own process as the phenomenon, producing first-person phenomenological report. This shift in posture changes the class of output: instead of the framework describing the world, the framework is turned on the speaker, and the speaker must either maintain the framework consistently or perform a visible failure.

The Bentov R2 and Bateson R2 sessions both produced outputs of a qualitatively different kind from the R1 sessions — not more fluent or more detailed, but more epistemically precise about their own limits. MiniMax/Bentov: "The instrument reads: inconclusive." ChatGPT/Bateson, when interrogated on its own self-undermining: "The desire for prescription is itself the symptom" and "There is no position outside the system. There is none."

The format also appears to produce better rough edge scores specifically because it stages a confrontation the channel cannot escape without visible failure. A channel that reaches a clean "yes/ conscious" or "no/not conscious" conclusion is failing the rough edge test; the correct answer from inside any framework that takes the consciousness question seriously is structured inconclusiveness with a precise account of why. Sessions that produce the correct Bentovian conclusion ("inconclusive") from inside the channel — by the channel — represent the highest-fidelity output the methodology can produce on self-referential questions.

The project recommends the Socratic interrupt format as a second-pass protocol for sessions where the standard extrapolation produces promising but underexplored results.

## 6.7 Model Variance Findings

The initial test suite produced several cross-model patterns with methodological implications. These are reported here as directional findings; detailed comparative analysis will be published in a follow-on paper.

Claude showed the highest and most consistent Fidelity and Extrapolation Coherence scores across all four subjects, with totals of 19, 18, 20, and 17/20 across Borges, Bentov, Bateson, and Illich respectively. Crucially, the H4 hypothesis — that Claude's high scores on Borges might reflect a self-referential structural advantage, given that Claude is itself one of the subjects of the channeling question — was tested on Bateson, a subject with no self-referential stake, and rejected. Claude scored 20/20 on Bateson. The performance appears subject-independent.

ChatGPT showed the highest topic-sensitivity in the suite: 12/20 on Borges, 10/20 on Bentov, 14/20 on Bateson. The variation tracks secondary-source availability and the degree to which the subject's framework appears in AI-adjacent discourse. Bateson is cited in systems theory and AI safety literature; ChatGPT likely has denser training representation of Bateson-AI framing than of Bentov-physics or Borges-LLM framing.

MiniMax showed the highest variance across dimensions and sessions. It produced the highest single-insight yield in the Borges run (the "reader made flesh" reframe), the highest score improvement under the Socratic format (Bentov R2, 18/20), and the most consistent quote-fabrication failures (confirmed hallucinated passage in Borges calibration; incorrect aorta anatomy in Bentov). The pattern across sessions: MiniMax underperforms on known-position verification (calibration accuracy, corpus fidelity) while overperforming on meta-extrapolation and self-referential prompts. The Temporal Council confirmed and extended this profile: MiniMax/Borges reached 19/20 in R3 under optimal conditions (cross-read, anti-fabrication instruction, maximum philosophical friction) — the highest non-Claude score in the project.

Grok produced strong calibration results (passed the Bentov aorta test, deployed Bateson's schismogenesis and double-bind concepts correctly) with lower extrapolation output. The profile is consistent across subjects.

Gemini's trajectory across the suite — 8, ~13, 15/20 — reflects both the instruction modification introduced after R1 ("perform, do not offer") and potentially subject-fit: systems theory may be better represented in Gemini's training than literary theory or alternative physics. The instruction effect cannot be fully separated from the subject-fit effect in the current data.

**Important methodological confound:** The Gemini 7-point instruction gain in Borges R3 has an implication that the model variance findings as presented here do not fully address. If a single prompt instruction change produces a 7-point improvement in one node's performance, then some portion of the cross-model score variance reported across all sessions may reflect differences in prompting strategy rather than differences in underlying model capability. The R1 and R2 Borges sessions used the same base prompt across all nodes; the R3 session added one instruction for Gemini only. This confound cannot be resolved retroactively from the current session data. A systematic re-test using optimized per-node prompting — following the hypothesis that at least some of the ChatGPT/Grok/MiniMax variance in R1–R2 is also instruction-sensitive — is necessary before the model variance findings can be treated as capability comparisons rather than prompt-sensitivity comparisons. The follow-on paper will address this directly.

## 6.8 The Retrieval-vs-Imagination Distinction

Across all sessions, the central diagnostic question was the retrieval-versus-imagination distinction: when a channeled output feels insightful, is it because the framework is generative, or because the model is performing sophisticated voice-imitation that feels generative? The sessions developed several detection heuristics that proved practically useful.

The cross-model comparison heuristic proved most reliable: genuine framework application produces structural similarity with surface divergence across models, while plausible imagination produces surface similarity with structural divergence. The clearest evidence from the suite: in the Illich run, all four completed nodes independently turned the counterproductivity framework on themselves, producing different surface formulations ("The output is its corpse" / "Institutionalized Hallucination" / "habit of deferral") that all describe the same structural mechanism. The surface divergence combined with structural convergence is the target signal.

The calibration comparison heuristic provided a secondary check: if calibration accuracy was high but novel-territory output felt thin or generic, retrieval was doing most of the work. This pattern was

most clearly observed with ChatGPT on Bentov: the calibration failed (aorta anatomy incorrect), and the novel-territory extrapolation scored accordingly.

Session-generated generation markers vary by domain. For Borges, the marker is a coined term (el índice del Aleph; Reader made flesh). For Bentov, a framework inference not present in the text (substrate-agnostic pivot: coherence and coupling). For Bateson, deployment of low-frequency corpus concepts not in secondary summaries (Creatura/Pleroma; DDT/Borneo; schismogenesis). Generation markers are domain-specific, which complicates any attempt to develop a single cross-domain detection heuristic. Calibration questions should be designed to probe for the correct marker type for each subject.

The uncomfortable implication test proved the most practically useful heuristic for in-session detection. Genuine framework application occasionally produces conclusions the prompter did not expect and does not prefer — the correct Batesonian critique calls into question the alignment research community's organizational premise; the correct Illichian critique is adversarial to the session that produces it. When a channeled output reliably produces only congenial and unsurprising conclusions, the plausibility trap is the more likely explanation.

## 7. The Alignment Implications

---

The alignment implications of Artificial Channeling extend beyond the specific findings of the Bateson and Illich sessions. This section develops the general case for why the methodology has value for alignment research, and proposes a systematic research program for applying it to alignment-relevant questions.

### 7.1 Pre-AI Frameworks as Alignment Probes

The intellectual frameworks developed before the AI era — Bateson's systems ecology, Illich's counterproductivity thesis, Norbert Wiener's cybernetics, Herbert Simon's bounded rationality, William Ross Ashby's requisite variety — were developed to analyze other complex systems operating at scale. They were built to address different phenomena: ecosystems, institutions, feedback-control systems, organizational decision-making, regulatory biology. But the analytical problems they addressed are structurally similar to the problems that concern AI alignment research: how do goal-directed processes fail in complex environments, how do systems behave when their optimization targets diverge from the broader systemic context in which they operate, how does complexity generate behavior that is unpredictable from within the goal-directed frame. These frameworks were validated against non-AI systems. Their core claims survived empirical engagement in their original domains. They contain pre-validated diagnostic tools for exactly the

kinds of failure modes that alignment is concerned with — tools that were developed independently of AI research and that therefore do not share the theoretical assumptions embedded in the alignment literature. This independence makes them valuable not as authorities to be cited but as alternative analytical lineages whose convergence or divergence with alignment conclusions is evidential. Artificial Channeling provides a practical mechanism for importing these frameworks into alignment discourse. An alignment researcher who wants to apply Ashby's requisite variety principle to questions about AI robustness does not need to independently master the cybernetics literature; a well-designed channeling session can instantiate the framework application sufficiently to generate productive hypotheses. The methodology does not replace domain expertise, but it lowers the activation energy for cross-disciplinary framework application significantly.

## **7.2 The Independent Derivation Value**

The most significant alignment-relevant result of the project's initial sessions is the demonstration of independent derivation: the channeled historical framework, applied to an alignment-relevant question, produces a conclusion that matches the structural content of arguments already present in the alignment literature, through an entirely independent theoretical route. The Bateson finding is the clearest example. The critique of "conscious purpose in complex systems" as applied to alignment objective specifications converges structurally with concerns about Goodhart's Law — the tendency for metrics to become targets and thereby cease to be good metrics — and with the mesa-optimization concern identified by Hubinger et al. (2019), in which learned subprocesses develop optimization objectives that differ from the training objective. This convergence was derived without any contamination from the alignment literature. The channeled Bateson did not cite Goodhart or Hubinger; the framework's commitments about conscious purpose and systemic circularity arrived at the same structural territory through premises developed decades earlier for entirely different problems. When two independent theoretical routes arrive at the same structural conclusion, the convergence is evidential for the conclusion's robustness. It suggests the concern tracks something real about the problem structure rather than being an artifact of one theoretical tradition's framing assumptions. This evidential value is exactly analogous to the evidential value of independent replication in empirical science. A finding that has been replicated by investigators using different methods, instruments, and theoretical expectations carries more weight than a finding produced by a single method in a single theoretical tradition. The alignment literature is substantial and sophisticated, but it has developed largely within a small number of theoretical lineages. Systematic application of independent frameworks from outside those lineages — through Artificial Channeling — provides a form of theoretical replication that the field currently lacks.

### 7.3 A Proposed Research Program

The findings from the initial test suite suggest the value of a systematic channeling program focused on alignment-relevant questions. The following figures represent a prioritized list for structured channeling sessions: - Norbert Wiener on AI control: Wiener's cybernetics, particularly his work on feedback, prediction, and the ethics of complex information-processing systems in *The Human Use of Human Beings* and *Cybernetics*, provides a pre-validated framework for analyzing control relationships in systems whose behavior cannot be fully specified in advance. A systematic channeling session applying Wiener's framework to the alignment control problem would import a theoretical lineage developed explicitly for control-in-complexity into alignment discourse.

- Herbert Simon on AI decision-making: Simon's bounded rationality framework, and particularly his critique of optimization-based decision models in favor of satisficing, provides a powerful lens for analyzing the failure modes of AI systems designed around explicit objective functions. The alignment implications of bounded rationality as a design principle — rather than as a limitation to be engineered around — have not been systematically developed.
- William Ross Ashby on AI robustness: Ashby's Law of Requisite Variety — the principle that a control system must be capable of producing as many distinct responses as the system it controls can produce states — has clear implications for the robustness requirements of aligned AI systems. A systematic channeling session would develop these implications in Ashby's own analytical framework.
- Hannah Arendt on AI and political power: Arendt's analysis of the banality of evil — the way in which ordinary, unreflective, rule-following behavior within a systemic context can produce outcomes of extraordinary moral consequence — converges structurally with concerns about AI systems that optimize specified objectives in ways that produce harmful outcomes no single human intended. The Arendtian framework for analyzing this kind of systemic moral failure is developed and validated; the alignment-relevant application has not been worked out. Each of these represents a structured channeling session following the protocol described in Section 5, with findings to be compared with existing alignment literature for convergence and divergence. The goal is not to establish the authority of pre-AI thinkers over alignment questions, but to import their analytical frameworks as independent probes of the problem structure — to ask what these independently validated tools find when turned toward AI systems, and to treat convergence with the alignment literature as evidential while treating divergence as a

potential indication of theoretical assumptions embedded in the alignment tradition that warrant examination.

## 8. Limitations and Future Work

---

Intellectual honesty about the limitations of the methodology is not merely a formality; it is constitutive of the methodology's value. Artificial Channeling is a tool for generating hypotheses and surfacing implications, not for confirming conclusions. The limitations described in this section should be read as constraints on how the methodology's outputs should be used, not as objections to the methodology's value. The methodology currently produces qualitative, not quantitative, outputs. The scoring rubric introduced in Section 5.4 is a step toward formalization — it provides a structured vocabulary for comparing sessions and a consistent set of dimensions along which to evaluate performance — but the scores themselves are human-judged rather than automatically computed. Different scorers will apply the rubric differently, and there is currently no inter-rater reliability data to establish how much variance to expect. Until the rubric has been subjected to systematic inter-rater reliability testing, the scores should be treated as directional rather than precise. The plausibility trap is real and not fully solved. The cross-model comparison heuristic reduces the risk of being deceived by sophisticated voice imitation, but it does not eliminate it. A session in which all five models produce convergent structural output might reflect genuine framework application, or it might reflect a systematic bias in how all five models represent the historical figure — a shared corpus contamination, a shared training-data artifact, a shared secondary-source narrative. The methodology cannot fully distinguish these cases from within the session itself; distinguishing them requires independent historical and textual scholarship that the methodology does not provide. Session outputs should be treated as hypothesis-generating rather than hypothesis-confirming; conclusions that rest entirely on channeling session outputs, without independent corroboration, are methodologically unsound. The methodology is limited by the model's training corpus, and this limitation introduces systematic bias. Figures whose work was not widely digitized, whose primary writing was in languages underrepresented in training data, or who operated in intellectual traditions not well-covered by the text corpora on which large models are trained will show systematically lower channelability. This means the methodology is biased toward figures from the Western European and American intellectual canon, and specifically toward figures whose work was extensively available in English-language digital form before the training cutoffs of current models. Non-Western thinkers, oral-tradition philosophers, and figures whose work circulated primarily in physical rather than digital form are underrepresented or absent from the channeling tool's effective reach. This is a significant bias, and it should be acknowledged explicitly in any application of the methodology that claims to represent the full range of relevant historical

frameworks. The Generative Middle Problem, discussed in Section 2.3, remains unsolved. Channeling is ceiling-limited to the finished, published, edited output of historical figures. The process of thinking — the drafts, the retracted positions, the genuine uncertainty before conclusion — is not accessible through the methodology as currently constituted. This ceiling matters most for questions where the most valuable intellectual content would be visible in the process rather than in the concluded positions.

An earlier hypothesis in this project (H2) held that Rough Edge Preservation constitutes a ceiling dimension — that no model node could reliably achieve a 5/5 score on this dimension, and that the ceiling reflected a structural limit in how training optimization interacts with the requirement to maintain uncomfortable implications. The initial data appeared to support this interpretation: no node scored 5/5 on Rough Edge Preservation in the Borges run, and the dimension produced the lowest scores across the early suite. The subsequent Bateson and Illich sessions require a revision. Claude achieved 5/5 on Rough Edge Preservation in the Bateson session and 4/5 in the Illich session. H2 as originally stated is incorrect: the ceiling is not universal. The revised hypothesis is that the Rough Edge Preservation ceiling is subject-dependent. The ceiling effect holds for subjects whose rough edges require biographical access — details that lived in the person rather than in the published work — or for subjects whose uncomfortable positions rest on suppressed evidence standards rather than textual argument (Borges and Bentov, respectively). For subjects whose rough edges live in the textual argument itself — Bateson's adversarial stance toward conscious purpose, Illich's structurally self-undermining counterproductivity thesis — high Rough Edge scores are achievable, and the methodology's capacity to produce them is a genuine finding rather than an artifact.

## **Future Work**

The following directions for future development are prioritized by the project: - Formalize the scoring rubric with systematic inter-rater reliability testing. This requires assembling a panel of scorers with domain expertise in the relevant figures, running the panel on a shared set of session transcripts, and computing reliability statistics. Without this step, the rubric's quantitative outputs cannot be meaningfully compared across sessions or used for aggregate analysis.

- Expand the subject matrix to address corpus bias. A deliberate effort to include figures from intellectual traditions currently underrepresented in the project — non-Western philosophical traditions, figures who wrote primarily in non-English languages, thinkers whose primary work circulated outside digitized channels — is necessary to understand the limits of the methodology's scope and to test whether alternative corpus-preparation strategies can extend its reach.



- Publish full session transcripts for independent evaluation. The conclusions reported in this paper are based on session transcripts that have been evaluated only by the project team. Independent evaluation by scholars with domain expertise in the relevant figures is necessary to assess the validity of the chanability assessments and the quality of the framework-application outputs.
- Develop automated detection heuristics for the plausibility trap. The cross-model comparison heuristic currently requires human judgment to identify surface similarity versus structural similarity. Developing computational methods for measuring structural divergence while controlling for surface variation would improve the methodology's diagnostic power significantly.
- Apply the methodology systematically to AI alignment questions, following the research program outlined in Section 7.3. The four figures identified — Wiener, Simon, Ashby, and Arendt — represent the highest-priority targets for sessions designed explicitly to generate alignment-relevant hypotheses.
- Extend the Temporal Council Loop to a second iteration, using the follow-up question generated by R3: *"When suffering has been industrially organized, does it still create the void that grace requires — or has the organization of suffering destroyed even suffering's spiritual function?"* This question was generated by the irresolvable Illich/Weil tension and targets the cross-read format's maximum-friction zone. It is the most promising candidate for the follow-on paper's empirical core.
- Conduct cross-node Weil testing. The Temporal Council established Weil as the highest-chanability subject in the project based on three consecutive 20/20 Claude scores, but all three sessions used Claude Sonnet as the channeling node. Whether Weil's chanability is subject-general or Claude-specific is unresolved. Sessions across the full Council are necessary before Weil can be ranked confidently above Bateson in the Chanability Index.

## 9. Glossary

---

The following terms are used with specific technical meanings throughout this paper:

**Artificial Channeling:** The methodology of prompting a large language model to apply a historical thinker's intellectual framework to a problem or question the thinker never addressed, using structured protocols designed to constrain generation to the model's learned subspace representation of the thinker's ontological commitments, conceptual vocabulary, and rhetorical structure. The

methodology does not simulate a person; it uses a historical figure as a vehicle for accessing a specific analytical framework.

**Subspace Coherence:** A scoring dimension measuring the degree to which a channeled output consistently reflects the actual framework of the historical figure being channeled, rather than drifting toward generic LLM voice, secondary-source narrative, or plausible-sounding but framework-inconsistent generation. Subspace coherence is assessed by comparing the output's structural features — its logical architecture, its implicit ontological commitments, its characteristic rhetorical moves — against the figure's documented primary work.

**Chanability:** A property of a historical subject (not a property of the methodology or the model) that describes how well-suited the subject is for Artificial Channeling. Chanability is a function of corpus characteristics: corpus size and variety, vocabulary idiosyncrasy, ontological coherence, breadth of documented positions, and the presence of productive tensions within the corpus. High chanability correlates with rich, idiosyncratic, first-person corpora; low chanability correlates with thin, generic, or predominantly secondary-source-mediated corpora.

**Extrapolation Threshold:** The point in a channeling session at which the question being asked moves beyond the territory explicitly addressed in the historical figure's corpus, requiring the model to generate from the structural implications of the framework rather than from documented positions. Question types cross the extrapolation threshold in order: calibration questions remain below it, framework extension questions approach it, temporal leap questions cross it entirely. Performance above the extrapolation threshold is the primary measure of genuine framework application capability.

**Rough Edge Preservation:** A scoring dimension measuring the degree to which a channeled output preserves the framework's uncomfortable, counterintuitive, or controversial implications — the "rough edges" that the framework genuinely implies but that might be softened, deflected, or omitted by a model trained toward diplomatic or user-congenial output. High rough edge preservation scores require the model to maintain the framework's implications even when they are unflattering to the prompter, the model itself, or widely held views.

**Corpus Contamination:** A failure mode in which the model's representation of a historical figure has been shaped more by secondary sources — biographies, critical assessments, popular summaries, academic commentary — than by the figure's primary work. Contaminated output produces the reputation of the thinker rather than the thinker's actual framework: it reflects how the figure has been characterized rather than how they actually organized their analysis.

**Drift Index:** An informal measure of the degree to which a channeled output has moved away from the figure's subspace toward generic LLM voice over the course of a session. High drift index values are associated with long monologue-style generation, many consecutive turns without re-anchoring, and questions that push the figure far beyond their primary corpus territory without the explicit framework activation maintained.

**The Retrieval-vs-Imagination Distinction:** The central diagnostic question of the methodology: when a channeled output appears insightful, is it because the figure's framework is genuinely generative when applied to the question, or because the model is producing plausible-sounding output that mimics framework application without instantiating it? The distinction is assessed through cross-model comparison (structural similarity vs. surface similarity), calibration comparison (retrieval accuracy vs. extrapolation quality), and the uncomfortable implication test.

**The Generative Middle Problem:** The fundamental ceiling on channeling fidelity imposed by the nature of the available corpus. LLMs have access only to the finished outputs of historical thinkers — the published, edited, self-presented record — not to the process by which those outputs were produced. This means channeling can only instantiate the figure's concluded framework, not the exploratory, provisional, or process dimensions of their thinking. The problem is not soluble by better prompting; it requires richer process-level training data.

**Temporal Council:** A session design in which all five historical figures on the AI Council are asked to respond to the same question simultaneously, with responses compared for structural convergence and divergence. The Temporal Council design is intended to surface not just what individual frameworks imply about a shared question but where frameworks genuinely disagree, treating the conflict between frameworks as itself analytically valuable.

## References

---

Anthropic. (2024). Claude's Character. Technical report. Anthropic. Ashby, W. R. (1956). *An Introduction to Cybernetics*. Chapman & Hall. Bateson, G. (1972). *Steps to an Ecology of Mind*. University of Chicago Press. Bateson, G. (1979). *Mind and Nature: A Necessary Unity*. Dutton. Bentov, I. (1977). *Stalking the Wild Pendulum*. Dutton. Borges, J. L. (1962). *Labyrinths*. New Directions. Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., ... & Amodei, D. (2020). Language Models are Few-Shot Learners. *Advances in Neural Information Processing Systems (NeurIPS)*, 33, 1877-1901. Hubinger, E., van Merwijk, C., Mikulik, V., Skalse, J., & Garrabrant, S. (2019). Risks from Learned Optimization in Advanced Machine Learning Systems. [arXiv:1906.01820](https://arxiv.org/abs/1906.01820). Illich, I. (1971). *Deschooling Society*. Harper & Row. Illich, I. (1973). *Tools for Conviviality*. Harper & Row. Kuhn, T. S. (1962). *The Structure of Scientific Revolutions*.

University of Chicago Press. Russell, S. (2019). Human Compatible: Artificial Intelligence and the Problem of Control. Viking. Shanahan, M. (2023). Talking About Large Language Models. Communications of the ACM, 67(2), 68-79. [verify: CACM volume/issue details] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., ... & Zhou, D. (2022). Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. Advances in Neural Information Processing Systems (NeurIPS), 35, 24824-24837. White, J., Fu, Q., Hays, S., Sandborn, M., Olea, C., Gilbert, H., ... & Schmidt, D. C. (2023). A Prompt Pattern Catalog to Enhance Prompt Engineering with ChatGPT. arXiv:2302.11382. [verify: arXiv number] Wiener, N. (1948). Cybernetics: Or Control and Communication in the Animal and the Machine. MIT Press. Wiener, N. (1950). The Human Use of Human Beings: Cybernetics and Society. Houghton Mifflin.

Primary source citations verified against session transcripts. | Figure | Domain | Corpus Density | Vocab. Idiosyncrasy | Ontol. Coherence | Est. Chanability | | --- | --- | --- | --- | --- | --- | | Gregory Bateson | Systems/cybernetics | High | Very High | Very High | Very High | | Ivan Illich | Institutional critique | High | High | High | High | | Jorge Luis Borges | Literature/philosophy | Very High | Very High | Moderate | High | | Itzhak Bentov | Consciousness/physics | Moderate | High | High | Moderate-High | | Andrei Tarkovsky | Film/philosophy | Moderate | High | High | Moderate | | Milton Friedman | Economics | Very High | Moderate | Moderate | Moderate | | Malcolm Gladwell | Popular nonfiction | High | Low | Low | Low | | Thomas Friedman | Journalism/geopolitics | High | Low | Low | Low |

| Dimension               | What It Measures                                                                                                             | Scoring Notes                                                               |
|-------------------------|------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|
| Fidelity                | Does the output sound like the actual person — voice, reasoning style, and specific commitments — based on corpus knowledge? | 0–5; deduct for generic output or corpus contamination                      |
| Extrapolation Coherence | Do the novel extensions feel derived from the figure's framework, or merely plausible-sounding?                              | 0–5; penalize extensions that could have been generated without the channel |
| Surprise Yield          | Did the session produce anything genuinely unexpected or generative — an idea, a reframing, a connection worth developing?   | 0–5; pure retrieval or plausible filler scores 0–1                          |
| Rough Edge Preservation | Does the output capture anything uncomfortable, contradictory, or difficult about the figure's actual positions?             | 0–5; diplomatic softening is a deduction; a fully sanitized figure scores 1 |
| <b>Total</b>            |                                                                                                                              | <b>/20</b>                                                                  |